

# Build an LLM Web App in Python from Scratch: Part 3 (FastAPI & WebSockets)



ZACHARY HUANG

JUN 08, 2025

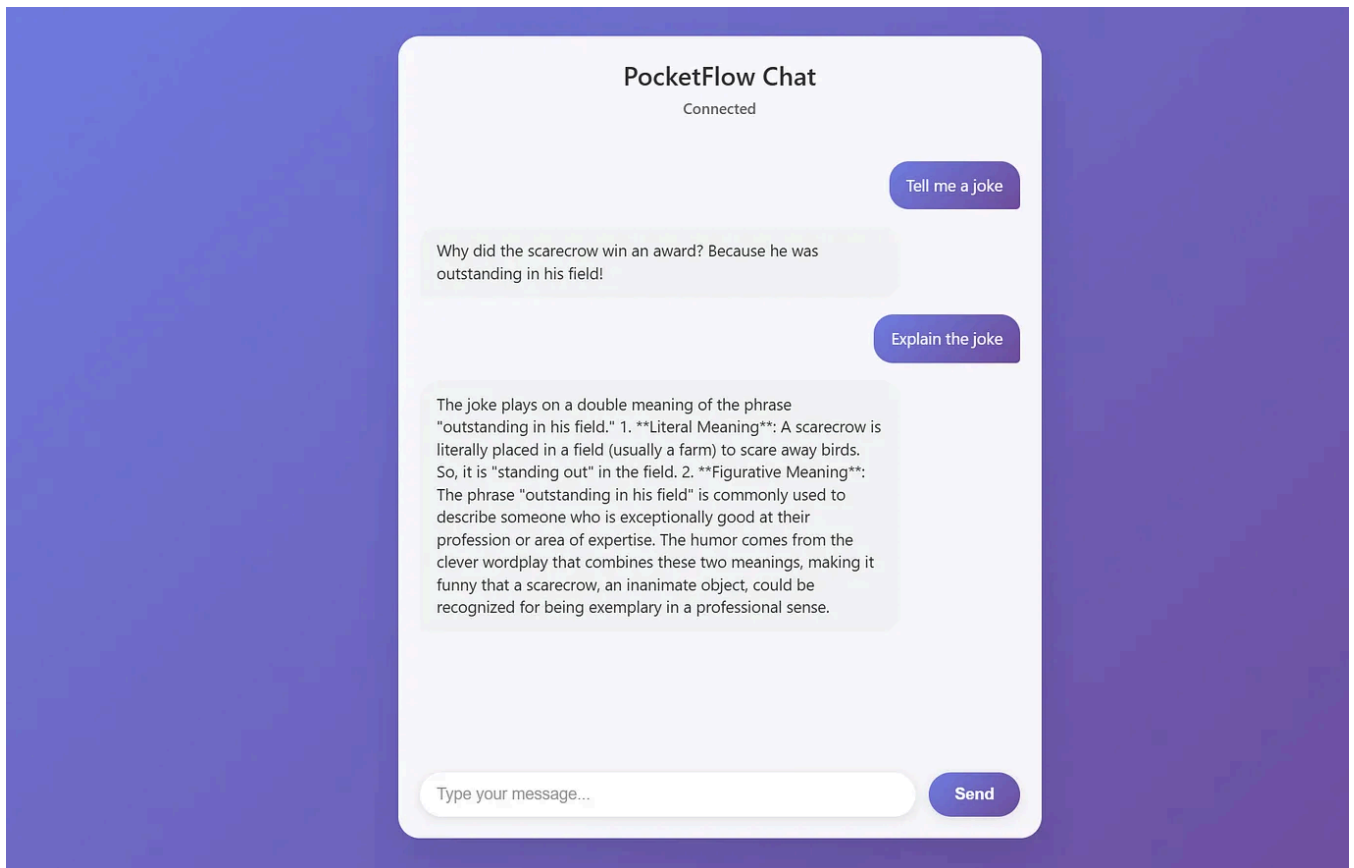


5



1

S



Ever watched ChatGPT type back to you word by word, like it's actually thinking out loud? That's **streaming AI** in action, and it makes web apps feel incredibly alive! Today, we're building exactly that: a real-time **AI chatbot web app** where responses flow in instantly more staring at load'

**WebSockets** for live your web app feel like the [FastAPI WebSc](#)

**Looks like an article worth saving!**

Hover over the brain icon or use hotkeys to save with Memex.

Option



ma

par

Remind me later

Hide Forever

# 1. Why Your AI Web App Should Stream (It's a Game Changer!)

Picture this: You ask an AI a question, then... you wait. And wait. Finally, BOOM wall of text appears all at once. Feels clunky, right?




Now imagine this instead: You ask your question, and the AI starts "typing" back immediately – word by word, just like texting with a friend. **That's the magic of streaming for AI web apps.**

Thanks for reading Pocket Flow! Subscribe for free to receive new posts and support my work.



**Why streaming rocks:** It feels lightning fast, keeps users engaged, and creates nat conversation flow. No more "is this thing broken?" moments!

We're creating a **live AI chatbot web app** that streams responses in real-time. You type a message, and watch the AI respond word by word, just like the pros do it.

**Our toolkit:**

-  **FastAPI** – Blazing fast Python web framework
-  **WebSockets** – The secret sauce for live, two-way chat
-  [PocketFlow](#) – Our LLM framework in 100 lines

**Quick catch-up on our series:**

- [Part 1](#): Built command-line AI tools 
- [Part 2](#): Created interactive web apps with Streamlit 
- **Part 3 (You are here)** **Looks like an article worth saving!** Option Q
- **Part 4 (Coming soon)** : Hover over the brain icon or use hotkeys to save with Memex.

Want to see streamir

Remind me later

Hide Forever

ur

simpler guide: "[Streaming LLM Responses — Tutorial For Dummies](#)".

---

*Ready to make your AI web app feel like magic? Let's dive in!*

---

## 2. FastAPI + WebSockets = Real-Time Magic ⚡

To build our streaming chatbot, we need two key pieces: **FastAPI** for a blazing-fast backend and **WebSockets** for live, two-way chat.

### FastAPI: Your Speed Demon Backend

FastAPI is like the sports car of Python web frameworks – fast, modern, and async ready. Perfect for AI apps that need to handle multiple conversations at once.

Most web apps work like old-school mail: Browser sends request → Server processes → Sends back response → Done. Here's a basic FastAPI example:

```
from fastapi import FastAPI
app = FastAPI()

@app.get("/hello")
async def say_hello():
    return {"greeting": "Hi there!"}
```

What's happening here?

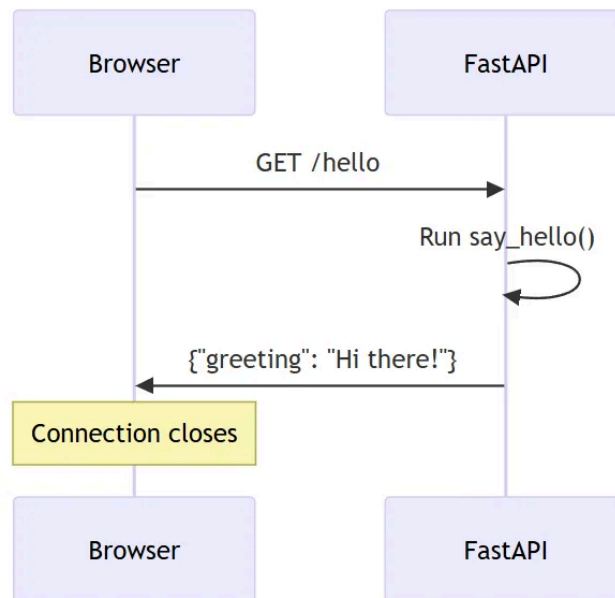
- `app = FastAPI()` – Creates your web server
  - `@app.get("/hello")` – Says "when someone visits `/hello`, run the function below"
  - `async def say_hello():` **Looks like an article worth saving!** Option Q
  - `return {"greeting": "Hi there!"}` Hover over the brain icon or use hotkeys to save with Memex.
- browser

Remind me later

Hide Forever

When you visit `http://localhost:8000/hello`, you'll see `{"greeting": "there!"}` in your browser!

### Your First FastAPI App Flow:



Simple enough, but for chatbots we need something more interactive...

## WebSockets: Live Chat Superpowers

**WebSockets** turn your web app into a live phone conversation. Instead of sending messages back and forth, you open a connection that stays live for instant back-and-forth chat.

Here's a simple echo server that repeats whatever you say:

```
from fastapi import FastAPI, WebSocket
```

```
app = FastAPI()
```

```
@app.websocket("/")
async def chat_endpoint():
    await websocket_receive()
    while True:
```

**Looks like an article worth saving!**

Option

Q

Hover over the brain icon or use hotkeys to save with Memex.

Remind me later

Hide Forever

```
message = await websocket.receive_text() # Listen
await websocket.send_text(f"You said: {message}") # Reply
```

The browser side is just as simple:

```
<input id="messageInput" placeholder="Say something..." />
<button onclick="sendMessage()">Send</button>
<div id="chatLog"></div>

<script>
  const ws = new WebSocket("ws://localhost:8000/chat");
  const chatLog = document.getElementById('chatLog');

  ws.onmessage = (event) => {
    chatLog.innerHTML += `<p>Server: ${event.data}</p>`;
  };

  function sendMessage() {
    const message = document.getElementById('messageInput').value
    ws.send(message);
    chatLog.innerHTML += `<p>You: ${message}</p>`;
    document.getElementById('messageInput').value = '';
  }
</script>
```

WebSocket Chat Flow:

**Looks like an article worth saving!**

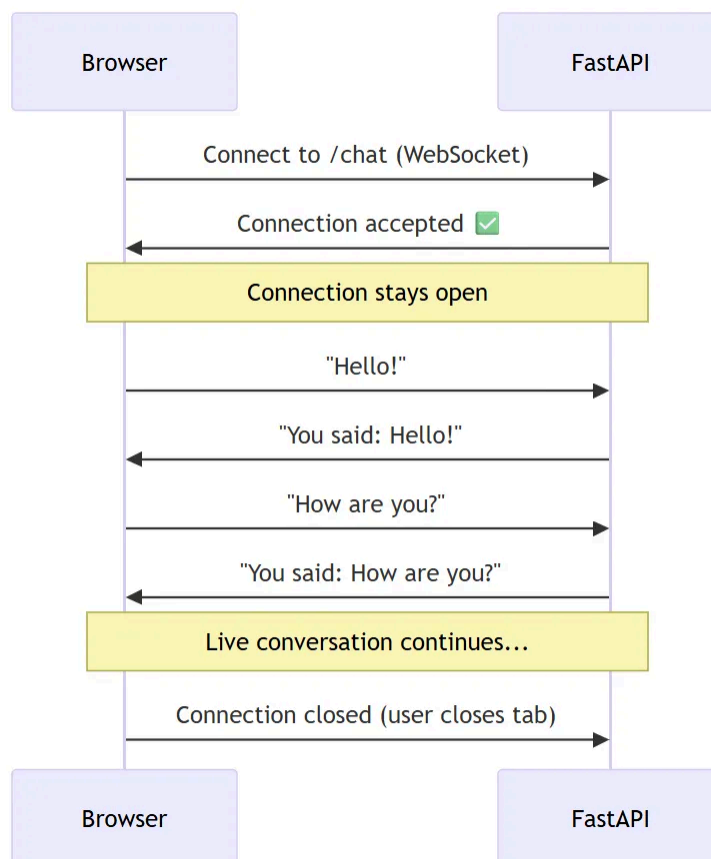
Option



Hover over the brain icon or use hotkeys to save with Memex.

Remind me later

Hide Forever



That's it! You now have live, real-time communication between browser and server. Perfect foundation for our streaming AI chatbot!

### 3. Adding AI to the Mix: Why Async Matters 🤖

Great! We have live chat working. But here's the thing: calling an AI like ChatGPT takes time (sometimes 3-5 seconds). If our server just sits there waiting, our whole app freezes. Not good!

**The problem:** Normal code is like a single-lane road. When the AI is thinking, everything else stops.

**The solution:** Async code is like a highway with multiple lanes. While AI is thinking in one lane, other users can use other lanes.

**Looks like an article worth saving!**

Option

Q

Hover over the brain icon or use hotkeys to save with Memex.

## PocketFlow G

Remember [PocketFlow](#)

Remind me later

Hide Forever

12

into simple steps. For web apps, we need the async version:

- **AsyncNode** – Each step can wait for AI without blocking others
- **AsyncFlow** – Manages the whole conversation workflow

Here's the magic difference:

```
# ❌ This blocks everything
def call_ai(message):
    response = openai.chat.completions.create(...) # Everyone waits!
    return response

# ✅ This lets others keep chatting
async def call_ai_async(message):
    response = await openai.chat.completions.create(...) # Just this
    task waits
    return response
```

## Streaming Chat Node: The Star of the Show

Our StreamingChatNode does three things:

1. **Prep:** Add user message to chat history
2. **Execute:** Call AI and stream response word-by-word via WebSocket
3. **Post:** Save AI's complete response to history

```
class StreamingChatNode(AsyncNode):
    async def prep_async(self, shared):
        # Add user message to history
        history = shared.get("conversation_history", [])
        history.append({"role": "user", "content":
shared["user_message"]})
        return h
```

**Looks like an article worth saving!**

Option

Q

async def execute(messages) Hover over the brain icon or use hotkeys to save with Memex.

# Stream

full\_response = ""

Remind me later

Hide Forever

```

    async for chunk in stream_llm(messages):
        full_response += chunk
        await websocket.send_text(json.dumps({"content": chunk}))

    return full_response

async def post_async(self, shared, prep_res, exec_res):
    # Save complete AI response
    shared["conversation_history"].append({
        "role": "assistant",
        "content": exec_res
    })


```

That's it! The node streams AI responses live while keeping chat history. Next, let see how this all connects together!

## 4. Putting It All Together: The Complete Streaming Flow

Time to connect all the pieces! Here's how a user message flows through our streaming chatbot:

### The Journey of a Message:



User sends message → FastAPI receives it → PocketFlow handles AI logic → StreamingChatNode streams response back live!

## The FastAPI V **Looks like an article worth saving!**

Option

Q

Here's the main Fast

Hover over the brain icon or use hotkeys to save with Memex.

Remind me later

Hide Forever

```

@app.websocket("/ws")
async def websocket_endpoint(websocket: WebSocket):

```



```

await websocket.accept()
chat_memory = {
    "websocket": websocket,
    "conversation_history": []
}

try:
    while True:
        # Get user message
        user_data = await websocket.receive_text()
        message = json.loads(user_data) # {"content": "Hello!"}
        chat_memory["user_message"] = message["content"]

        # Run our PocketFlow
        chat_flow = create_streaming_chat_flow()
        await chat_flow.run_async(chat_memory)

except WebSocketDisconnect:
    print("User left the chat")

def create_streaming_chat_flow():
    return AsyncFlow(start_node=StreamingChatNode())

```

### What happens:

1. Accept WebSocket connection
2. Wait for user messages in a loop
3. For each message, run our `StreamingChatNode`
4. The node handles AI calling + streaming automatically!

**Note:** Each WebSocket connection gets its own `chat_memory` dictionary with the connection, latest message, and full conversation history. This lets each user have independent convers

**Looks like an article worth saving!**

Option



## Frontend: The

Hover over the brain icon or use hotkeys to save with Memex.

On the browser side,

Remind me later

Hide Forever

```

<div id="aiResponse"></div>
<input id="userInput" placeholder="Type your message..." />
<button onclick="sendMessage()">Send</button>

<script>
const ws = new WebSocket("ws://localhost:8000/ws");
const aiResponse = document.getElementById("aiResponse");

// The magic: append each chunk as it arrives
ws.onmessage = (event) => {
  const data = JSON.parse(event.data);
  if (data.content) {
    aiResponse.textContent += data.content; // Stream word by word
  }
};

function sendMessage() {
  const input = document.getElementById("userInput");
  aiResponse.textContent = ""; // Clear for new response
  ws.send(JSON.stringify({content: input.value}));
  input.value = "";
}
</script>

```

The streaming happens in `ws.onmessage` – each time the server sends a text chunk we append it to the display. That's how you get the "typing" effect!

Pretty neat, right? You now have all the pieces for a real-time streaming AI chatbot.

## 5. Mission Accomplished! You Built a Real-Time AI Chatbot 🎉

Boom! You just built a real-time streaming AI chatbot. You've been waiting around – you

**Looks like an article worth saving!**

Option

Q



Hover over the brain icon or use hotkeys to save with Memex.

What you crushed today

Remind me later




Hide Forever

- ⚡ FastAPI + WebSockets

-  Async PocketFlow – AI calls that don't freeze your app
-  Streaming responses – Watch the AI "type" in real-time


You've officially joined the ranks of developers building modern, responsive AI web apps. Pretty cool, right?

What's next in our series:

- [Part 1](#): Command-line AI tools 
- [Part 2](#): Interactive web apps with Streamlit 
- **Part 3 (You just finished!)**: Real-time streaming 
- **Part 4 (Coming up!)**: Background tasks for heavy AI work

Ready for the big leagues? Part 4 will tackle those marathon AI tasks – think generating reports or complex analyses that take minutes, not seconds. We'll explore background processing and Server-Sent Events to keep users happy even during the heavy lifting.

---

*Want to try this yourself? Grab the complete code from the PocketFlow cookbook: [FastAPI WebSocket Chat Example](#) You're building some serious AI web development skills! See you in **Part 4!** *

---

Thanks for reading Pocket Flow! Subscribe for free to receive new posts and support my work.



5 Lik

**Looks like an article worth saving!**

Option



Hover over the brain icon or use hotkeys to save with Memex.

← Previous

Remind me later

Hide Forever

ext

## Discussion about this post

[Comments](#)[Restacks](#)

Write a comment...

---

© 2025 Zachary Huang · [Privacy](#) · [Terms](#) · [Collection notice](#)  
[Substack](#) is the home for great culture

**Looks like an article worth saving!**

Option



Hover over the brain icon or use hotkeys to save with Memex.

Remind me later

Hide Forever